Bayesian Linear Regression

Summary

Given:

A set of training points $\{(x_i, y_i)\}_{i=1}^N$, where $x_i \in \mathbb{R}^d$ and $y_i \in \mathbb{R}$

Generative model for the data in terms of the weight vector w and hyperparameters, as described below.

Find:

Query 1: The posterior probability distribution of w.

Metrics:

Metric 1: Expected squared Euclidean distance between the predicted mean $\hat{w}$ and the true mean w, where the expectation is taken with respect to the posterior distribution.

Metric 2: Total variation distance between the computed posterior and the correct posterior over w.

Details

The file "problem-1-generator.R" contains R code to generate the true regression coefficients and the input training data. The model is

$$
\begin{aligned}
\Sigma_1 &= 2\mathbf{I}_{5\times5} \\
\mu &\sim \mathcal{N}(0, \Sigma_1) \\
\Sigma_2 &= \mathbf{I}_{5\times5} \\
\Sigma_{\text{prior}} &\sim \text{Wishart}(1, \Sigma_2) \\
w &\sim \mathcal{N}\left(\mu, \Sigma_{\text{prior}}^{-1}\right) \\
x_{ij} &\sim \text{Uniform}(-1, 1) \\
\tau &\sim \text{Gamma}(0.5, 2) \\
\epsilon_i &\sim \mathcal{N}\left(0, \frac{1}{\tau}\right) \\
y_i &= \sum_j x_{ij} w_j + \epsilon_i
\end{aligned}
$$

The file contains 500 training examples generated from a single run of the R code. There are four covariates generated uniformly from . The values of the variables that generated the data are

$$
\begin{aligned}
\mu &= (-1.8195312, 1.2237587, 0.8361809, -2.6017006, -2.3574193) \\
\Sigma_{\text{prior}} &= (\text{see "problem-1-prior.Sigma.csv"}) \\
w &= (-1.731855, 2.986017, 2.698284, -3.591651, -3.714157) \\
(x_i, y_i) &= (\text{see "problem-1-dat.csv"})
\end{aligned}
$$

Queries/Metrics:

1. Let $P(\hat{w}|D)$ be the posterior distribution of the estimated weight vector. One metric is the expected squared error $\mathbb{E}\left[\|w \hat{-} w\|^2\right]$ under this distribution.

2. We have provided samples generated from the true posterior distribution $P_{\text{true}}(\hat{w}|D)$. We can estimate the total variation distance between the true distribution and your estimate $P(\hat{w}|D)$ using the samples generated by your estimated distribution:

$$
\int_w \left| P\left(w\hat{|}D - P_{\text{true}}(\hat{w}|D)\right) \right| dw
$$