

Small Problem 1: Bayesian Linear Regression

Query 1: Posterior distribution of the estimated weight vector.

Metrics

Metric 1

Let $P(\hat{w}|D)$ be the posterior distribution of the estimated weight vector. Metric 1 is the expected squared error $\mathbb{E}[\|\hat{w} - w\|^2]$ under this distribution.

Metric 2

We have provided samples from the true posterior distribution $P_{true}(\hat{w}|D)$. Estimate the total variation distance between the provided samples and the samples from your estimate $P(\hat{w}|D)$:

$$\int_w |P(\hat{w}|D) - P_{true}(\hat{w}|D)| dw$$

Ground Truth:

For metric 1, the ground truth is the weight vector given in the problem details document. For ease of reference, this weight vector is also given here:

$$w = (-1.731855, 2.986017, 2.698284, -3.591651, -3.714157)$$

For metric 2, samples from the true posterior distribution are in the file `problem-1-solution-samples.xlsx` (without headers in the .csv version) and the columns represent the values of w_1, w_2, w_3, w_4, w_5 .

TODO:

Compute metric 1 using the ground truth weight vector given above.

Compute metric 2, the total variation between ground truth samples and the samples generated by your solution.

For the metric 2 computation, we have provided a Java program that computes the total variation between ground truth samples and the samples generated by your solution. The Java program is provided as both a jar file and as source code.

The jar file is `"totalvar.jar"`

The source code is located in folder `problem-1-total-variation-java`.

The file containing your samples should be a comma separated value (CSV) file without column headers where the columns represent the dimension of the weight vector (in this case there should be 5 columns corresponding to w_1, w_2, w_3, w_4 , and w_5), and the rows represent samples.

To evaluate your samples using the jar file, use the command line given below:

```
java -jar totalvar.jar <ground-truth-file-path> <input-file-path>
```

Evaluation output will be written to `stdout`.

To evaluate your samples using the provided source code, follow these steps:

Find the variable `static String dataFileA` in the `Main.java` file and set its value to be the complete path of the CSV file that contains the ground truth samples, e.g. the complete path to the `problem-1-solution-samples.csv` file referenced above.

Find the variable `static String dataFileB` in the `Main.java` file and set its value to be the complete path of the CSV file that contains your samples.

Run the program using the main method in this same `Main.java` file and the program will produce a total variation score. The method `computeScore` is also provided if you wish to programmatically evaluate your samples.

Note that we have included a detailed description of the “total variation” computation procedure in `problem-1-total-variation.pdf`.

Submission:

The metric value should be computed for each elapsed time step (by calling the provided code or by implementing yourself). The metric value should be reported for several elapsed time steps. The number of elapsed time steps should be sufficient to establish an “informative profile”.

For further details regarding submission of the metric and your code, please refer to the main CP4 problem description document, e.g. `PPAML-Challenge-Problem-4.pdf`.

Sample output files for this problem have been provided in the `sampleoutput` folder:

```
problem-1-query-1-metric-1.csv
problem-1-query-1-metric-2.csv
```

Ground Truth Details:

This problem can be solved using the [Stan probabilistic programming language](#) (Stan Development Team, 2014). Here is the Stan/R program that estimates the posterior distribution of the estimated weight vector.

Step 1: Save the following code as a Stan file (attached as `cp4_1.stan`)

```
data {
  int<lower=0> N;
  int<lower=1> K;
  matrix[N, K] x;
  vector[N] y;
}

parameters {
```

```

vector[K] w_mean;
cov_matrix[K] w_prec;
vector[K] w;
real<lower=0> noise_sd;
}

model {
  // Just the identity matrix
  matrix[K,K] idK;
  idK <- diag_matrix(rep_vector(1.0,K));

  w_mean ~ normal(0,2);
  w_prec ~ wishart(K, idK);
  //w ~ multi_normal_prec(w_mean, w_prec);
  w ~ normal(w_mean,1);
  // The second parameter is the scale 0.5,
  // corresponding to a rate parameter of 2.
  noise_sd ~ inv_gamma(0.5, 0.5);

  y ~ normal(x * w, noise_sd);
}

```

Step 2: Save the following code as a R file (attached as “cp4_1.R”)

```
setwd("~/galois/ppaml/challenge-problems/cp4/")

xy <- read.csv("problem-1-data.csv")
x <- xy[,2:6]
y <- xy[,7]

theData <- list(
  x=x
  , y=y
  , K=ncol(x)
  , N=nrow(x)
)

library(rstan)
sm <- stan_model("cp4_1.stan")

# For MAP estimate
map <- optimizing(sm, data=theData)

# To sample from posterior
fit <- sampling(sm, data=theData)

samp <- extract(fit)
write.table(samp$w, "cp4_1_w_posterior.csv", row.names=F, col.names=F, sep=',')
```